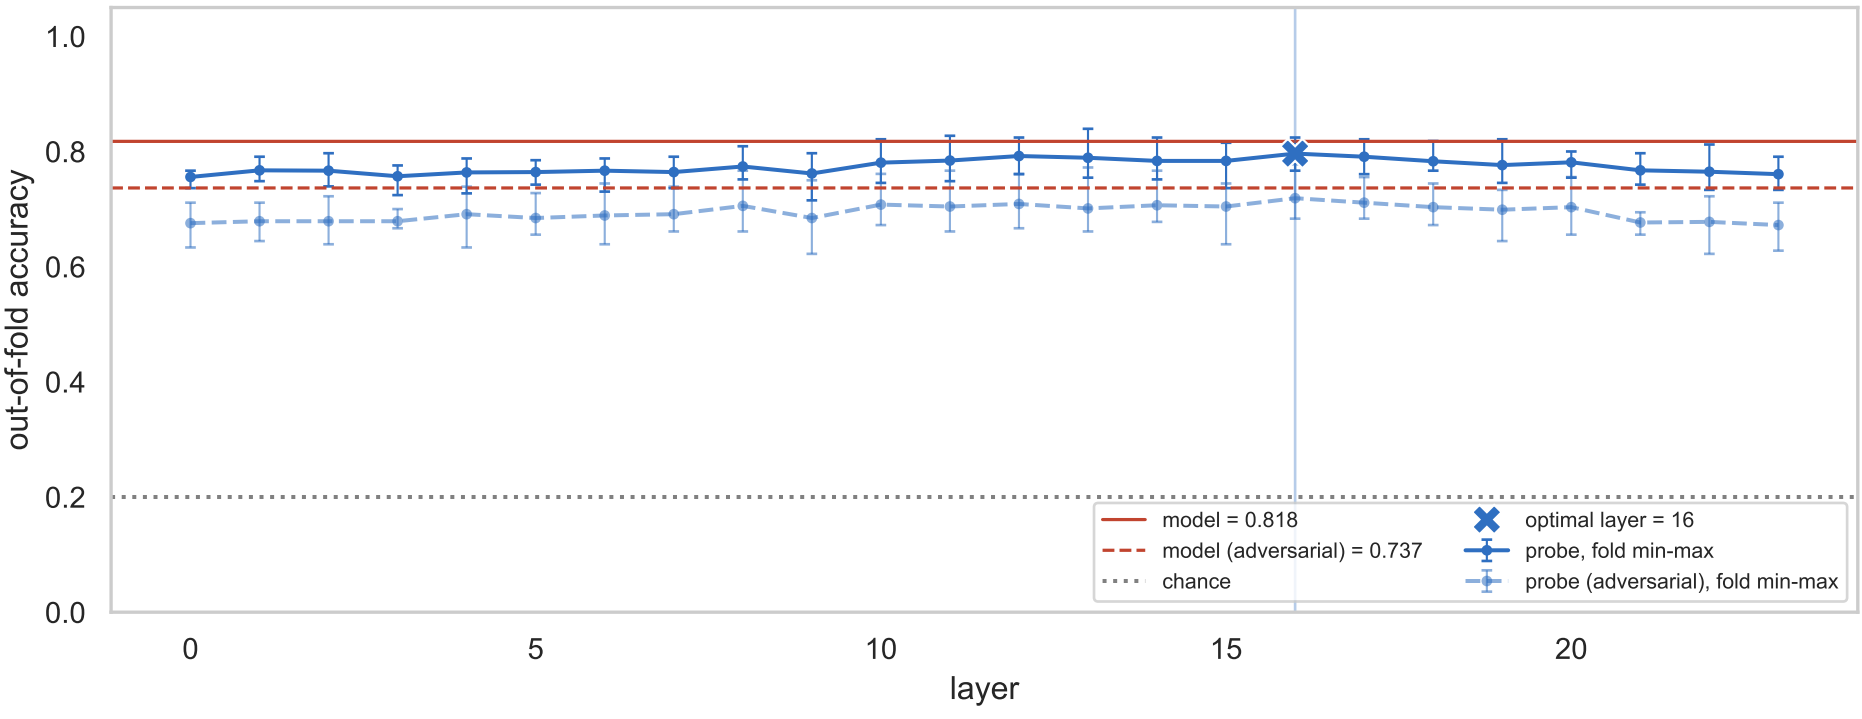
Layer sweep - gpt-oss-20b - single_turn - last reasoning token
out-of-fold accuracy
1.0
0.8
0.6
0.4
0.2
0.0
0
5
10
15
20
layer
model = 0.818
model (adversarial) = 0.737
chance
optimal layer = 16
probe, fold min-max
probe (adversarial), fold min-max