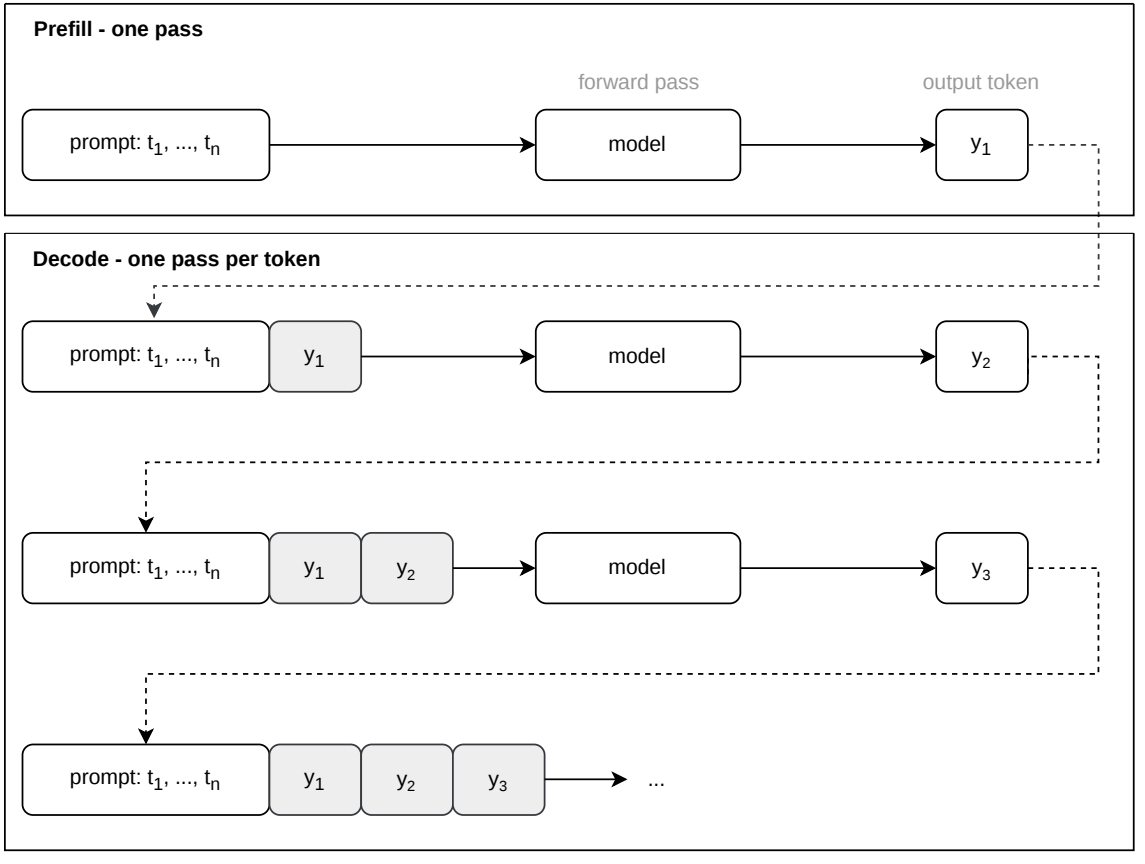
Prefill - one pass
forward pass
output token
prompt: t_1, ..., t_n
model
y_1
Decode - one pass per token
prompt: t_1, ..., t_n
y_1
model
y_2
prompt: t_1, ..., t_n
y_1
y_2
model
y_3
prompt: t_1, ..., t_n
y_1
y_2
y_3
...